

PolyChristoffel Networks

Obtaining latent manifolds via curvature learning

Mihir Talati

December 19, 2025

Motivation & Intuition

Why another dynamics model?

- Most current NNs are considered Black-Box simulators.
- We know something useful hides in Latent Space (see AlexNet) but we can't extract it due to arbitrary nonlinearities.
- Modern Models solve this by imposing structure on embeddings and evolutions.
- Physics models such as HNNs/LNNs do this via:
 - invariances, conservation laws, interpretable coefficients
 - extrapolation beyond the training distribution
- However we suspect that the geometry of latent spaces encode these properties

Geometric intuition: “forces” as coordinate effects and nonlinearity

- A straight-line path in one coordinate system can look curved in another.
- Example: inertial motion in Cartesian becomes nontrivial in polar coordinates.
- The apparent “forces” come from geometry via **Christoffel symbols**.

$$\ddot{x}^i + \Gamma_{jk}^i(x) \dot{x}^j \dot{x}^k = 0$$

- Interpret $\Gamma_{jk}^i(x)$ as **state-dependent couplings** of velocities into acceleration.
- Goal here: make the **equations of motion** a first-class object:
 - learn $\Gamma(x)$ so that trajectories follow a geodesic equation.

Graduate DiffGeom in 30 seconds: metric, connection, geodesics

- A (pseudo-)Riemannian metric $g(x)$ defines local inner products:

$$\langle u, v \rangle_x = u^\top g(x) v$$

- The Levi–Civita connection (torsion-free, metric-compatible) yields Christoffels:

$$\Gamma^i_{jk}(g) = \tfrac{1}{2} g^{i\ell} \left(\partial_j g_{k\ell} + \partial_k g_{j\ell} - \partial_\ell g_{jk} \right)$$

- Geodesics are “straightest” paths under that connection:

$$\ddot{x}^i + \Gamma^i_{jk}(x) \dot{x}^j \dot{x}^k = 0.$$

Model

Model overview

Inputs: initial state/latent x_0 , initial velocity v_0 (and optionally observed trajectory).

Learn: a polynomial Christoffel field $\Gamma(x)$ (and optionally an autoencoder).

Observed y_t	Latent/state x_t
$x_t = E(y_t) \Rightarrow$	$\Gamma(x)$ produces geodesic rollout $\hat{x}_{0:T}$
$\hat{y}_t = D(\hat{x}_t)$	Losses enforce dynamics + metric consistency

Design principle: keep the dynamics layer interpretable by making $\Gamma(x)$ a low-degree polynomial.

Parameterization: polynomial Christoffel network

Let $x \in \mathbb{R}^d$. Build monomial features $\phi(x) \in \mathbb{R}^P$ (degree $\leq p$):

$$\phi(x) = [1, x_1, \dots, x_d, x_1^2, x_1x_2, \dots].$$

Then define

$$\Gamma_{jk}^i(x) = \sum_{m=1}^P C_{jk,m}^i \phi_m(x).$$

- $C_{jk,m}^i$ are trainable coefficients (directly interpretable).
- Enforce torsion-free symmetry by construction:

$$\Gamma_{jk}^i = \Gamma_{kj}^i.$$

Dynamics block: geodesic equation as a neural layer

Use first-order state $z = (x, v)$ with $v = \dot{x}$:

$$\dot{x} = v, \quad \dot{v}^i = -\Gamma_{jk}^i(x) v^j v^k.$$

This defines a deterministic vector field $f_\theta(z)$ where parameters are the polynomial coefficients.

Interpretability:

- acceleration is quadratic in velocity and structured by $\Gamma(x)$
- no arbitrary activation needed in the dynamics law

Integrator: Scuffed ResNet rollout (discrete geodesic flow)

A simple differentiable integrator (semi-implicit Euler):

$$v_{t+1} = v_t + \Delta t \, a(x_t, v_t), \quad x_{t+1} = x_t + \Delta t \, v_{t+1},$$

where

$$a^i(x, v) = -\Gamma_{jk}^i(x) v^j v^k.$$

- Equivalent to stacking residual blocks with shared parameters.
- Stable, CUDA-friendly, easy to backprop through long rollouts.
- Later swap-in: RK4 / differentiable ODE solvers if needed.

Optional: jointly learned autoencoder with manifold constraint

When observations live in y -space (images, fields, etc.), use an encoder/decoder:

$$x_t = E_\psi(y_t), \quad \hat{y}_t = D_\psi(\hat{x}_t).$$

To enforce latent validity on a chosen manifold \mathcal{M} :

$$x \leftarrow \Pi_{\mathcal{M}}(x)$$

(e.g., sphere projection, hyperboloid retraction).

- Joint training: geometry losses backprop through E_ψ .
- Encourages a chart where dynamics are “as geodesic as possible.”

Why metric reconstruction / pseudo-Riemannian regularization?

Learning $\Gamma(x)$ freely can fit trajectories but yield a connection that is not Levi–Civita of any metric.

To encourage a **(pseudo-)Riemannian interpretation**, impose constraints consistent with:

- torsion-free: $\Gamma^i_{jk} = \Gamma^i_{kj}$
- metric compatibility: $\nabla g = 0$
- fixed signature (p, q) , non-degeneracy ($\det g \neq 0$)

Metric reconstruction idea: integrate the compatibility equation

Metric compatibility in coordinates:

$$\partial_i g_{jk} = \Gamma_{ij}^\ell g_{\ell k} + \Gamma_{ik}^\ell g_{j\ell}.$$

Treat this as a (path) ODE along a curve $x(s)$ from a basepoint x_0 :

$$\frac{d}{ds} g_{jk}(s) = \left(\Gamma_{ij}^\ell(x(s)) g_{\ell k}(s) + \Gamma_{ik}^\ell(x(s)) g_{j\ell}(s) \right) \frac{dx^i}{ds}.$$

- Choose $g(x_0) = g_0$ with fixed signature; integrate to get $g(x)$.
- If Γ is not metric-realizable, the reconstructed $g(x)$ becomes **path-dependent**.

Consistency loss: penalize path dependence (loop loss)

Compute $g(x)$ via two different paths:

- Path A: straight line $x_0 \rightarrow x$
- Path B: two-segment $x_0 \rightarrow x_m \rightarrow x$

Then define:

$$\mathcal{L}_{\text{loop}} = \|g^{(A)}(x) - g^{(B)}(x)\|_F^2.$$

Interpretation:

- pushes $\Gamma(x)$ toward connections that preserve some bilinear form
- provides extra signal even if you only supervise trajectories

Pseudo-Riemannian validity losses (practical)

Given reconstructed $g(x)$:

- Symmetry:

$$\mathcal{L}_{\text{sym}} = \|g - g^\top\|_F^2$$

- Non-degeneracy barrier (avoid singular metric):

$$\mathcal{L}_{\text{det}} = \text{softplus}(\alpha - \log |\det g|)$$

- Fixed signature (p, q) via eigenvalue sign penalties:

$$\mathcal{L}_{\text{sig}} = \sum_{i=1}^d \text{softplus}(-\beta s_i \lambda_i(g)), \quad s_i \in \{+1, -1\}.$$

Overall objective and training loop

Primary fit: match trajectories (latent or decoded):

$$\mathcal{L}_{\text{traj}} = \frac{1}{T} \sum_t \|\hat{x}_t - x_t\|^2 \quad \text{or} \quad \mathcal{L}_{\text{recon}} = \frac{1}{T} \sum_t \|D(\hat{x}_t) - y_t\|^2.$$

Regularize geometry:

$$\mathcal{L} = \mathcal{L}_{\text{traj/recon}} + \lambda_{\text{loop}} \mathcal{L}_{\text{loop}} + \lambda_{\text{det}} \mathcal{L}_{\text{det}} + \lambda_{\text{sig}} \mathcal{L}_{\text{sig}} + \lambda \|C\|_2^2.$$

- Backprop through (i) rollout integrator and (ii) metric reconstruction.
- Coefficients C remain interpretable (polynomial terms in Γ).

Training: Backprop + Computational Graph + Pytorch

Forward pass (one minibatch)

- Initialize from data (or latent): (x_0, v_0) .
- Rollout dynamics for $t = 0, \dots, T - 1$ with step Δt :

$$v_{t+1} = v_t + \Delta t a(x_t, v_t), \quad x_{t+1} = x_t + \Delta t v_{t+1},$$

where

$$a^i(x_t, v_t) = -\Gamma_{jk}^i(x_t) v_t^j v_t^k, \quad \Gamma_{jk}^i(x) = \sum_{m=1}^P C_{jk,m}^i \phi_m(x).$$

- Compute loss on the trajectory (and optional geometry regularizers):

$$\mathcal{L} = \sum_{t=1}^T \|x_t - x_t^{\text{true}}\|^2 + \lambda \mathcal{L}_{\text{geom}}.$$

Autodiff view: the rollout is a differentiable computation graph. Gradients flow

$$\mathcal{L} \rightarrow (x_{1:T}, v_{1:T}) \rightarrow a(\cdot) \rightarrow \Gamma(x_t) \rightarrow C.$$

Gradient signal on Christoffel coefficients

Because Γ is **linear** in the coefficients $C_{jk,m}^i$, the gradient has a clean form.

At a single step t :

$$\Gamma_{jk}^i(x_t) = \sum_{m=1}^P C_{jk,m}^i \phi_m(x_t) \implies \frac{\partial \Gamma_{jk}^i(x_t)}{\partial C_{qr,m}^p} = \delta_p^i \delta_j^q \delta_k^r \phi_m(x_t).$$

Acceleration coupling:

$$a_t^i = -\Gamma_{jk}^i(x_t) v_t^j v_t^k \implies \frac{\partial a_t^i}{\partial C_{jk,m}^i} = -\phi_m(x_t) v_t^j v_t^k.$$

Chain rule through time (BPTT):

$$\frac{\partial \mathcal{L}}{\partial C_{jk,m}^i} = \sum_{t=0}^{T-1} \left\langle \frac{\partial \mathcal{L}}{\partial a_t^i}, -\phi_m(x_t) v_t^j v_t^k \right\rangle + \frac{\partial \mathcal{L}_{\text{geom}}}{\partial C_{jk,m}^i}.$$

Interpretation: each coefficient update is a weighted sum of *state features* $\phi_m(x_t)$ times *velocity interactions* $v_t^j v_t^k$, so learned geometry can be inspected term-by-term (e.g., which monomials drive which couplings).

Proof of Concept experiment: polar inertial motion

Use a synthetic benchmark with known Christoffels:

$$\Gamma_{\theta\theta}^r = -r, \quad \Gamma_{r\theta}^\theta = \Gamma_{\theta r}^\theta = \frac{1}{r}$$

(others = 0), derived from Euclidean plane in polar coordinates.

- Train $\Gamma(x)$ to reproduce trajectories.
- Use diagnostics to compare learned components to ground truth.
- Verify torsion-free symmetry and reduced path dependence in reconstructed g .

Diagnostics I: Learned Christoffels vs. ground truth

What we plot

Evaluate the learned connection on a grid (r, θ) and extract:

$$\Gamma_{\theta\theta}^r(r, \theta), \quad \Gamma_{r\theta}^\theta(r, \theta)$$

Compare to ground truth:

$$\Gamma_{\theta\theta}^r = -r, \quad \Gamma_{r\theta}^\theta = \frac{1}{r}.$$

How to interpret

- **θ -invariance:** GT depends only on r , so learned heatmaps should be *nearly constant across θ* .
- **Shape:** $\Gamma_{\theta\theta}^r$ should look like a linear ramp in r ; $\Gamma_{r\theta}^\theta$ should decay like $1/r$.
- **Torsion-free check:** $\Gamma_{r\theta}^\theta \approx \Gamma_{\theta r}^\theta$.

Christoffel Diagnostics

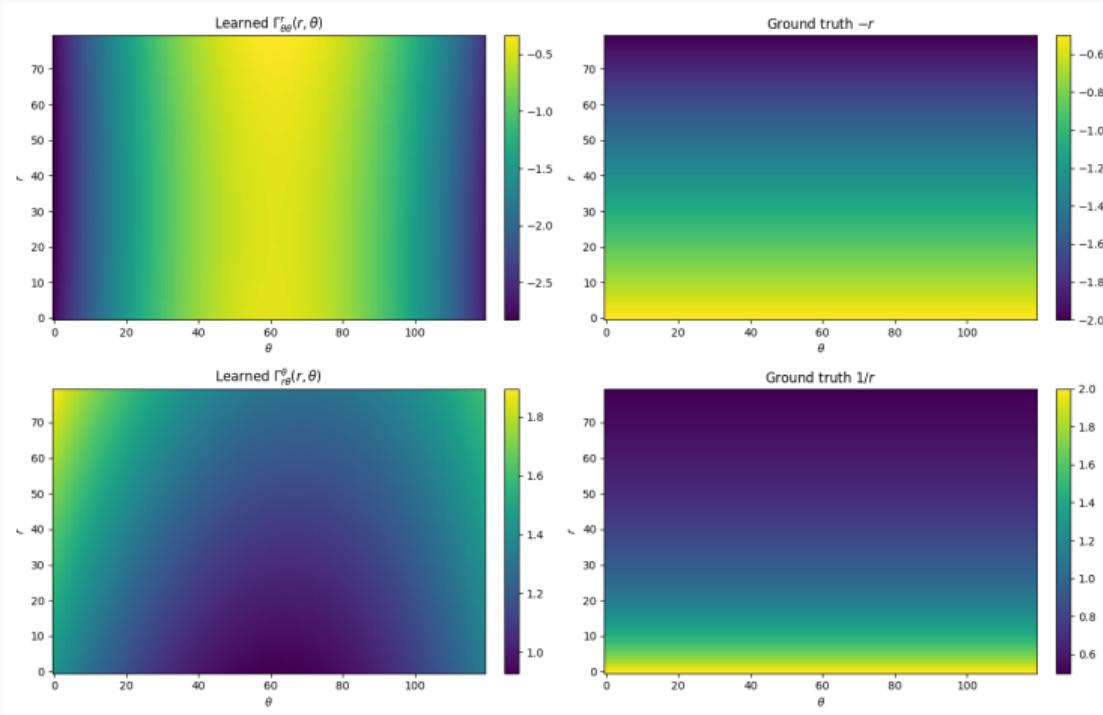


Figure 1: Learned vs. GT Christoffel components on (r, θ) grid.

(Christoffel Diagnostics w/AutoEncoder)

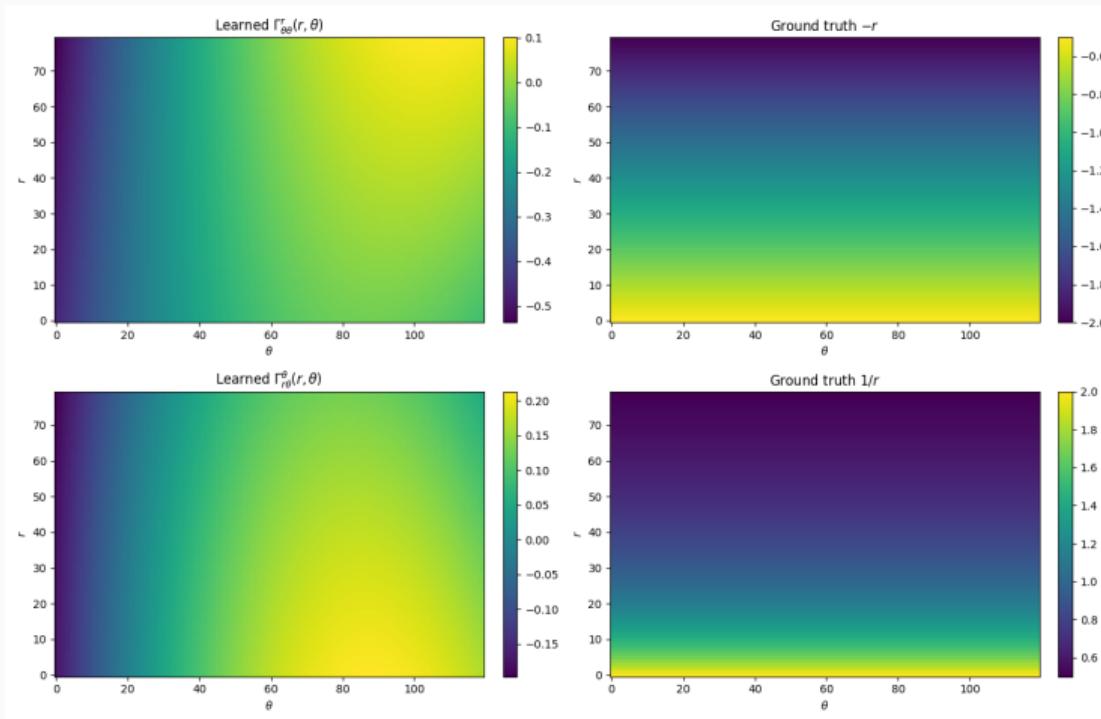


Figure 2: Learned vs. GT Christoffel components on **Embedded** (r, θ) grid.

Diagnostics II: Reconstructed metric and path-dependence (loop error)

Metric target (polar plane)

The Euclidean metric in polar coordinates is:

$$g(r, \theta) = \begin{pmatrix} 1 & 0 \\ 0 & r^2 \end{pmatrix}.$$

We reconstruct g from the learned Γ by integrating metric-compatibility:

$$\nabla g = 0 \quad \Rightarrow \quad \partial_i g_{jk} = \Gamma_{ij}^\ell g_{\ell k} + \Gamma_{ik}^\ell g_{j\ell}.$$

How to interpret

- **Diagonal structure:** $g_{r\theta} \approx 0$ (off-diagonal heatmap near zero).
- **Component shapes:** $g_{rr} \approx 1$ (flat), $g_{\theta\theta} \propto r^2$ (quadratic in r).
- **Path dependence:** reconstruct g via two paths; the loop error

$$\|g^{(A)}(x) - g^{(B)}(x)\|_F^2$$

should be small if Γ is close to Levi–Civita of some metric.

Pretty Graphs Pt 2

Reconstructed g components and loop error heatmap.

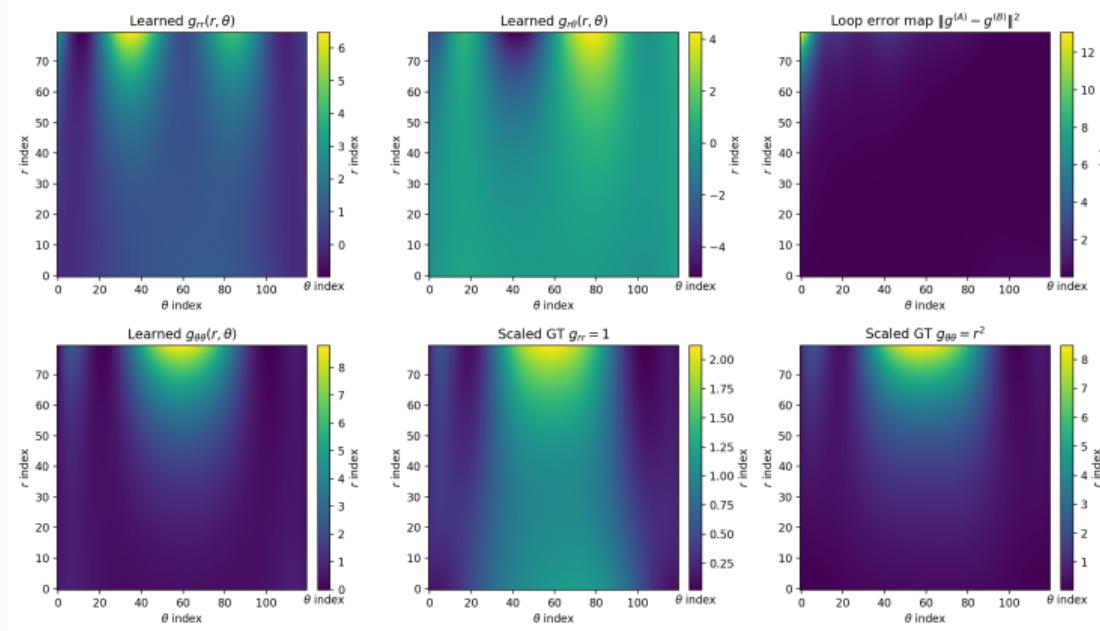


Figure 3: Learned Metric

Pretty Graphs Pt 2 Pt 2

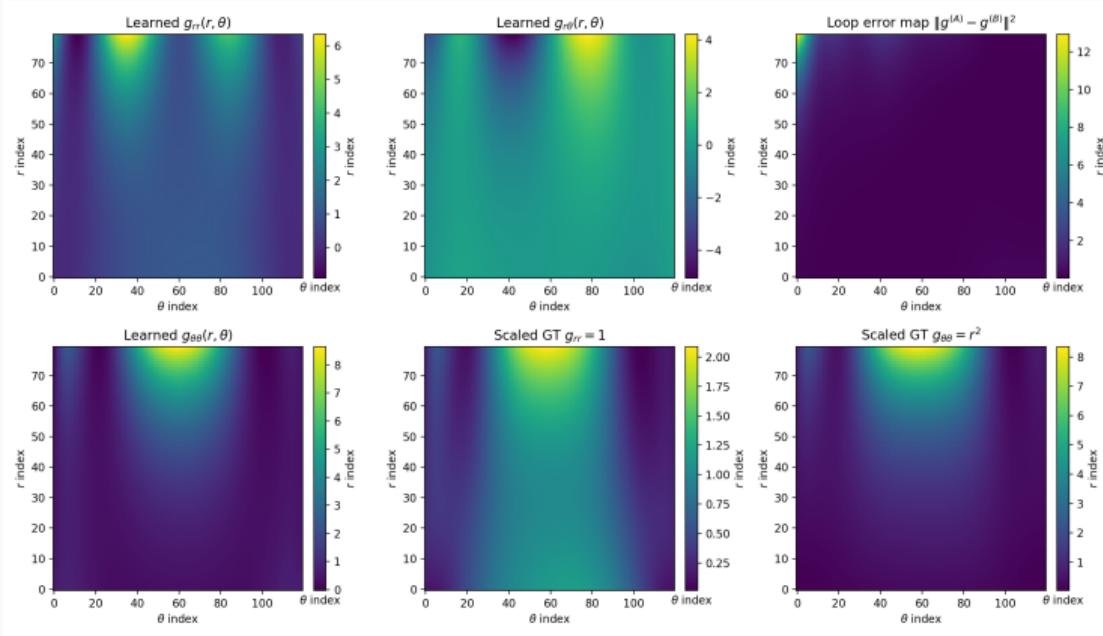


Figure 4: True Metric

AutoEncoded metrics)

While the AutoEncoded metrics are a little harder to interpret due to nonlinear coordinate transforms.

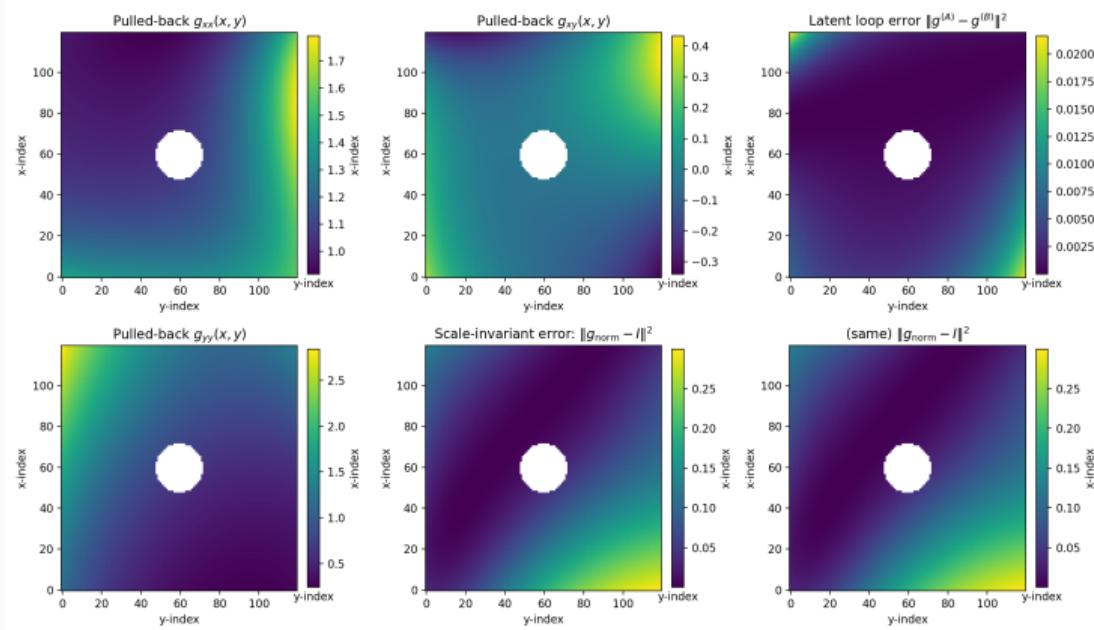


Figure 5: Learned Metric

AutoEncoded metrics)

We can still see what metric values we get by reinterpreting the embeddings under the ground truth metric.

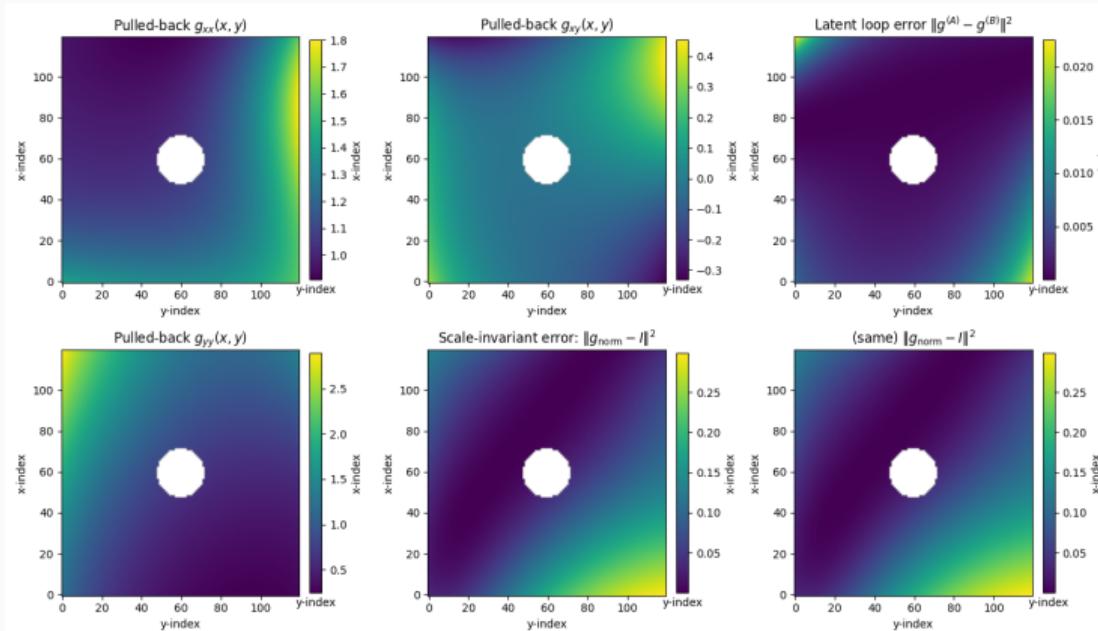


Figure 6: "True" *** Metric

Future Directions

Scaling up: leverage strong autoencoders (CNNs, etc.)

- Many domains already have excellent representation learners:
 - CNN autoencoders / VAEs for images and fields
 - pretrained encoders with semantic latent factors
- Proposal: **freeze or lightly finetune** an existing encoder/decoder, then learn $\Gamma(x)$ (and optionally $g(x)$) on the latent manifold to obtain:
 - interpretable latent dynamics
 - geometry-aware interpolation/extrapolation (geodesics)
 - constraints like signature / nondegeneracy for stability
- The notion of enforcing a Pseudo-Riemannian metric also defines a consistent, invariant dot product for vectors at any point on the manifold, which is also precisely how many encoders are trained (e.g in NLP)

Research directions

- **Metric Intervals:** Enforce a notion of "distance" via states through known equivalencies
- **Lorentzian Geometry:** Allow reparametrization of evolution along axes other than time
- **Beyond geodesics:** add forcing/dissipation in a structured way (e.g., geodesic + learned potential or Rayleigh dissipation term).
- **Better metrizability:** directly co-train $g(x)$ and enforce $\Gamma_{\text{poly}} \approx \Gamma(g)$ to guarantee Levi–Civita structure.
- **Topology/coverage:** multiple charts or implicit atlas, while retaining interpretability.
- **Stochastic extensions:** geodesic drift with learned diffusion (SDE in latent).

Takeaways

- Christoffel symbols provide an interpretable parameterization of dynamics:

$$\ddot{x}^i = -\Gamma_{jk}^i(x) \dot{x}^j \dot{x}^k$$

- Polynomial $\Gamma(x)$ yields a small, inspectable model with calculus-friendly smoothness.
- Metric reconstruction and pseudo-Riemannian regularizers provide extra training signals and geometric validity checks.
- Natural next step: apply on strong latent spaces (e.g., CNN-based autoencoders) to learn **curved latent dynamics**.

References i

-  T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. Duvenaud, "Neural Ordinary Differential Equations," arXiv:1806.07366, 2018.
-  S. Greydanus, M. Dzamba, and J. Yosinski, "Hamiltonian Neural Networks," arXiv:1906.01563, 2019.
-  E. O. Korman, "Atlas Based Representation and Metric Learning on Manifolds," arXiv:2106.07062, 2021.
-  J. Bürge, "Neural geodesic flows," Master's thesis, ETH Zürich, 2025.
-  N. He, M. Yang, and R. Ying, "Lorentzian Residual Neural Networks," arXiv:2412.14695, 2025.